

A Framework for Expert-Led Apps

Building **Lelanea.com** and a reusable framework for similar apps


Where we are now

Sunrise is our platform for building AI-driven apps, now in its first early release. One of our first projects on it is Lelanea, a transcendental coaching app which will be available at lelanea.com.

These apps **guide people through a personal growth journey** based on an expert's knowledge and experience. They should feel less like software and more like having a coach or a teacher by your side. Lelanea is the first example, but we're building a **framework**, not just a single app. Future apps will use the same foundation, making it easy to create learning and discovery experiences around any expert or body of work.

The five concepts

The framework rests on five concepts: **ways of thinking about the app** and of **organising its architecture**. They work together around the user:



1. Modules

Modules are the **places on the map**, or chapters of the expert's book, the parts of the app that the user naturally engages with. Each is self-contained, and each usually has its own **Chat Companion Window**: a conversation panel that helps and advises right where the user is working. The

framework ships no modules itself; instead, it defines **what a module is and how one is built**, and each app implements its own set of modules.

For Lelanea, the modules come straight from her “Becoming Whole” workbook and coaching specialities. Each stage of her journey becomes a place for the user to dwell and explore:


- **Core Values:** discovering what truly drives you.
- **Future Self:** picturing who you are becoming.
- **Lotus of Life:** a gentle look at how nourished each area of life feels.
- **Standards:** turning values into clear personal standards.
- **Boundaries:** mapping where boundaries are needed.
- **Liminality:** support through the in-between, threshold moments of change and transition.
- **Check-Ins:** short gratitude and reflection moments along the way.

Lelanea’s own body of work becomes the modules, the things the app learns, and the paths between them. The framework brings her work to life in a way that is personalised to every unique user and their life situations.

2. Facilitation Structures

This concept gives the experience **structure and direction**. As in a **map of towns and cities** joined by roads, some places open up only once others are visited, while some are free to explore any time. Facilitatory structures also keep the experience safe and consistent, providing good governance and watching over how the agents treat the user, and everything in it can be adjusted as we learn.

What this concept really provides are **routes**: enablement paths and enabling structures that open the way forward for the user. The point is not to constrain them to sequential steps but to **enable** them through many possible steps, laying down the paths along which they can grow and making the next worthwhile step possible.



For illustrative purposes only; not intended as a true reflection of Lelanea’s workbook.

The user feels this as gentle guidance: a warm welcome, check-ins, optional next steps, and a sense of where they have been and where they are heading.

A key job of this concept is to work out **what would serve the user's interests at any given moment**: whether to linger longer in the current module, or move on to another. It makes that judgement from the **freshest data-slots**, so the guidance responds to where the user actually is right now, not where they were when they started. The user will be invited to revisit and refine previous modules if the situation calls for it.

3. Data-Slots

A data slot is simply **one thing we want to know** about the user. Together, the slots build a rich picture that lets us personalise everything else.

The difference from an ordinary app: instead of filling in forms, **the user just talks**, and the system listens for what matters. Where it helps, we show them what we understood and let them correct or remove it. Some slots stay private to the system; most are open for the user to see and edit. If something comes up in a conversation that relates to a data slot outside of the subject matter of the conversation, it will be opportunistically noted and updated in the background so that nothing of importance gets lost, and we are continually building a broad understanding of the user and their life situations.

For Lelanea, the data-slots gather into natural areas of life, alongside deeper reads on the person:

- **Life areas**: physical health, emotional health, wealth, relationships, family, work, and spiritual health, including where things feel dysfunctional and where they are working well.
- **The person**: personality, life patterns, psyche, aspirations, and blind spots.
- **External conditions**: their current life situation and the constraints around them, so the app understands the circumstances they are actually living in.
- **Preferences**: smaller things, like how the user responds to certain kinds of questions, so the experience can adapt to them.

Data-slots also capture the **bespoke data tied to each module**. For Lelanea, that means the values, standards, boundaries and the rest are all held as slots, so the work the user does inside a module is captured and carried with them rather than lost when they move on.

Each data-slot answer is versioned and timestamped, so we keep a timeline of how the user evolves while using the app, building a rich picture that is never fixed and always changing with them. The system also attributes varying confidence levels to the data slots according to how the information was derived. If it is from a strongly worded comment directly from the user, it will be given strong confidence levels. If, however, the system inferred the data-slot answer from a tangential conversation, it will be given a weaker confidence score until firmed up through later user assertions.

Data-slots can receive confirmations and negations as conversations progress - firm data slots are further hardened when the user brings up something that confirms what we've already recorded. If the conversation contradicts a previously recorded slot, that is a negation which we can use as an opportunity to explore the reasons behind the contradiction - is the user changing their beliefs, or is it an area of uncertainty worth delving deeper into?

4. Knowledge Base

The Knowledge Base is the expert's own **intellectual property** (their books, talks, articles, methods) turned into a meaningful form that the AI can draw on. Their material is broken into pieces and stored by meaning, so that when a user asks something, the system can find the most relevant passages and put them in front of the AI before it answers.

This is what keeps the experience **authentic and trustworthy**. The AI speaks from the expert's actual work rather than from generic knowledge, which keeps it from making things up and gives the conversation rails to stay oriented around. For Lelanea, that means the app draws on her teaching, her podcast, and her workbook, so it sounds like her and stays true to her method.

Known in the industry as RAG (retrieval-augmented generation), the AI retrieves relevant material first, then answers from it. This works by a process called 'embedding' which attributes a range of numbers to words and phrases, storing them in a 'vector database'. Words with close meanings will be attributed similar numbers. Then, when a user types words and phrases into the chat, their sentences are also turned into numbers through an 'embedding' process, and these numbers can be compared to the numbers of the embedded content in the database to retrieve the expert's content that most closely matches.

5. Agents

Agents are AI personas, each with its own character and role. Some talk directly to the user; others work on processes in the background. We build and manage them; the user simply meets them through chat, a tap, voice, or even WhatsApp and email, and never sees the machinery.

Agents are what bring the other four concepts to life: they hold the conversations inside the **modules**, **draw on the Knowledge Base to stay grounded**, quietly fill and read the **data-slots**, and act on what the **Facilitation Structures** decide is the right next step.

They might be seen as a family of agents, or an **ecosystem of agents**, each playing a part and supporting the others. They can switch between empathy mode and teaching mode based on what the situation and the user's preferences call for (based on the data-slots storing user preferences around coaching styles).

For Lelanea, the family of agents might include:

- **Consultation & Discovery:** this agent builds a first picture of the person and their situation and shares it with the other agent, mirroring the person's own words back, so they feel heard.
- **Life Situation:** An agent specialised in probing into a certain life situation that the user presents. There could be sub-agents for this to help with Divorce, Job Loss, Relationship Struggles, Money Woes, Health Issues, etc.
- **Values, Transcending-Ego, Discernment, Boundaries & Liminality Agents:** each specialist agent is called in when working through a module or when the moment is right
- **Curator:** suggests a video, book or article from Lelanea's material.

- **Supervisor:** Oversees how the agents are interacting with the user to ensure the conversation is professional and adheres to the rules.

Illustrative only. The actual modules will emerge through designing Lelanea.

Agents rely on LLMs (Large Language Models) such as ChatGPT 5.x or Claude Opus, but heavily influence the 'reasoning' of these AI engines through the instructions we give them.

Agents have certain capabilities that can be configured by system admins. These capabilities could include having access to a certain area of the knowledge base, or the ability to do certain things within the app on behalf of the user, including clicking buttons, swiping screens or even searching the web. Agents can also be chained together to perform very complex tasks end-to-end - this is how a **Supervisor Agent** might monitor the behaviour of a **Coaching Agent**.

How the five concepts fit together

Each concept plays a distinct part in the overall experience:

- **Modules** are the places where the user spends time and does the work.
- **Facilitation Structures** lay down the routes between them: where the user can go and the range of next steps that are in the user's best interests.
- **Data-Slots** direct what the system learns about the user along the way.
- **Knowledge Base** is the expert's material that keeps it all grounded and authentic (prevents hallucinations).
- **Agents** are the voices and the hidden helpers that bring it to life.

The app meets the user wherever they are, across **desktop, iOS, Android, WhatsApp, SMS or Email**. If we get this right for Lelanea, and we earn a head start on every expert-led app that potentially follows.

An aspiration for the future:

Once the framework is mature, onboarding a new expert could become largely automatic. The domain expert uploads their knowledge base (their IP), and a dedicated agent, whose job is exactly this, designs the modules and facilitation paths, works out the data-slots worth capturing, and shapes the agent behaviour relevant to that domain. In other words, the structures that we design today by hand could one day be **self-generating** from the expert's material itself.

For now, the next step is to turn these five concepts into a working first version for Lelanea, and in doing so, lay the foundation that every later app could build on.